

BERT

Pre-training of Deep Bidirectional Transformers for Language Understanding

2019 @ Nations of the Americas Chapter of the Association for Computational Linguistics



URL to the
paper



AGENDA

04/11/2024

1 **The Big Picture:** BERT in the context of language modelling

2 A bird's eye view of the **architecture**

3 BERT's **versatility** – pre-training / fine-tuning process

4 **Conclusion & Discussion**

The (Single-Slide) Big Picture


BERT in the context of language modelling

THE BIG PICTURE: BERT IN THE CONTEXT OF LANGUAGE MODELING

Word2Vec is a static lookup table; BERT dynamically computes language encodings considering the whole sequence context.

Word2Vec: Static Lookup Table

🎯 Word2Vec aims to learn vector representations of words capturing their semantic / syntactic.



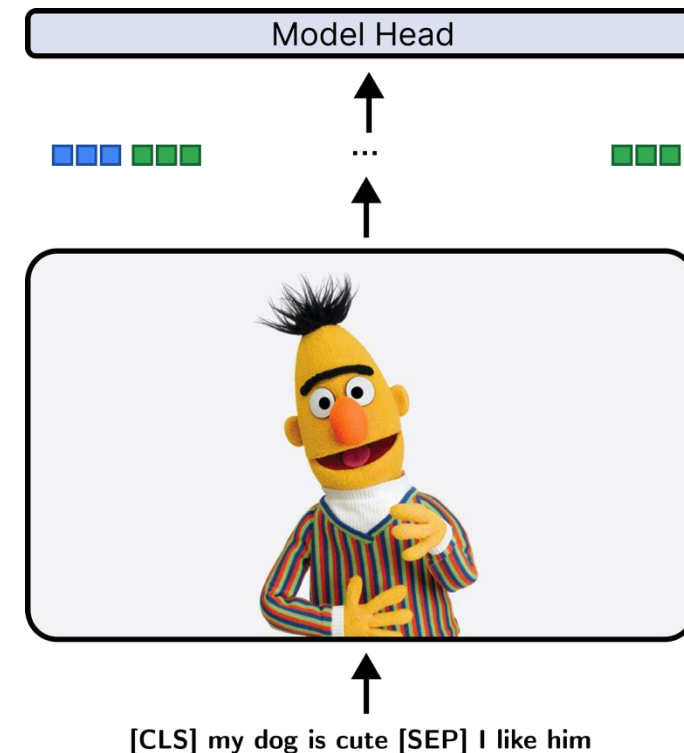
Has a fixed representation after the training ignoring current context of words.

*Please turn on the **light** vs. The feather is very **light***

BERT can solve different NLP tasks using different fine-tuned model heads.

BERT: Dynamic Language Understanding

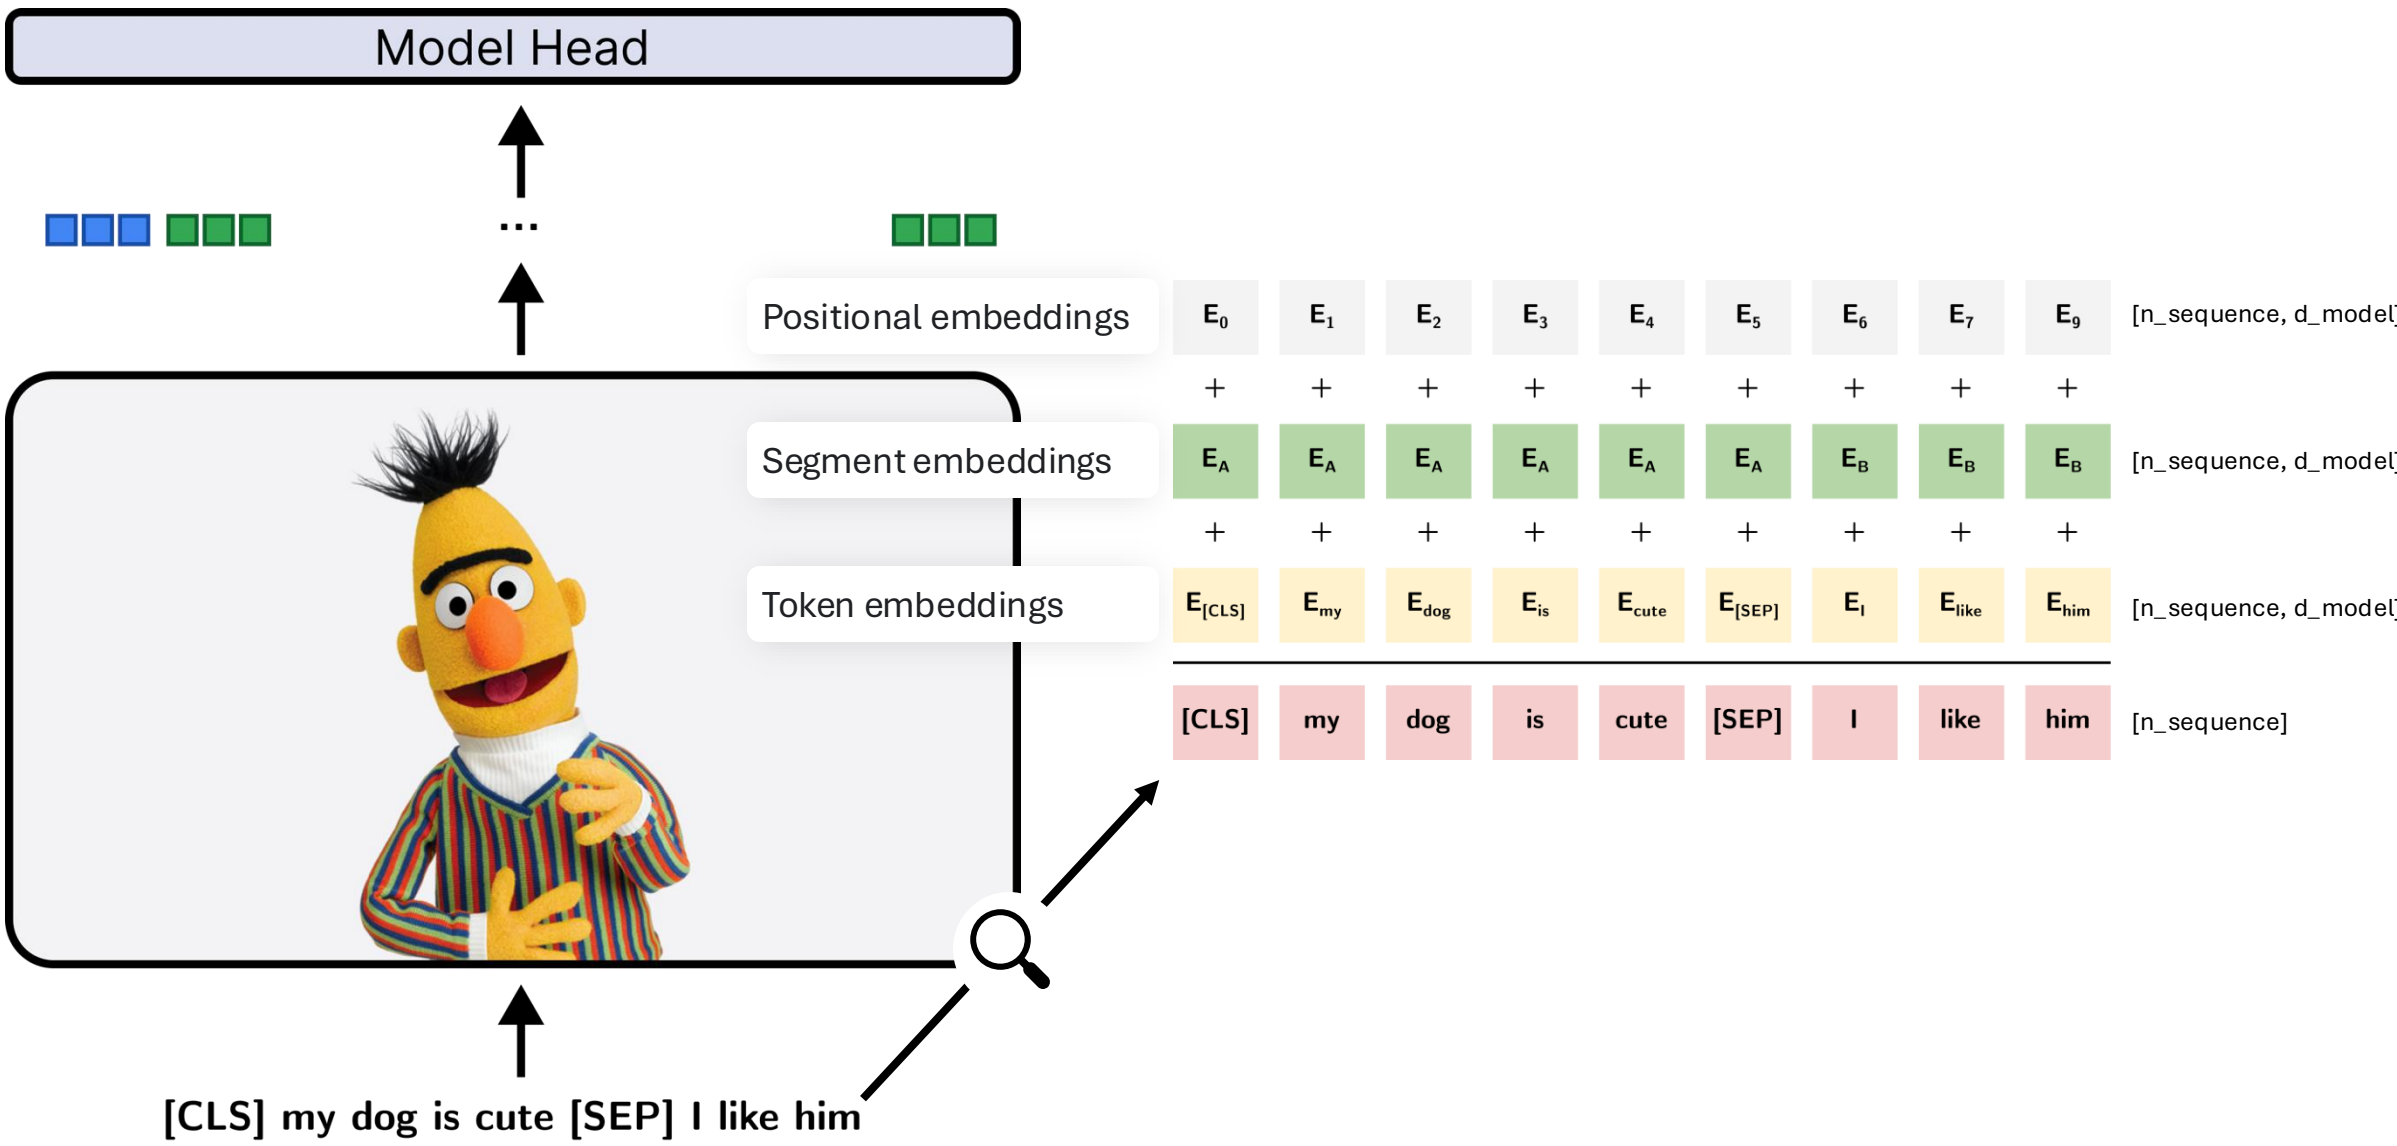
🎯 BERT aims to understand language by considering the full context of a word during training and inference.



A BERT's* eye view of the architecture.

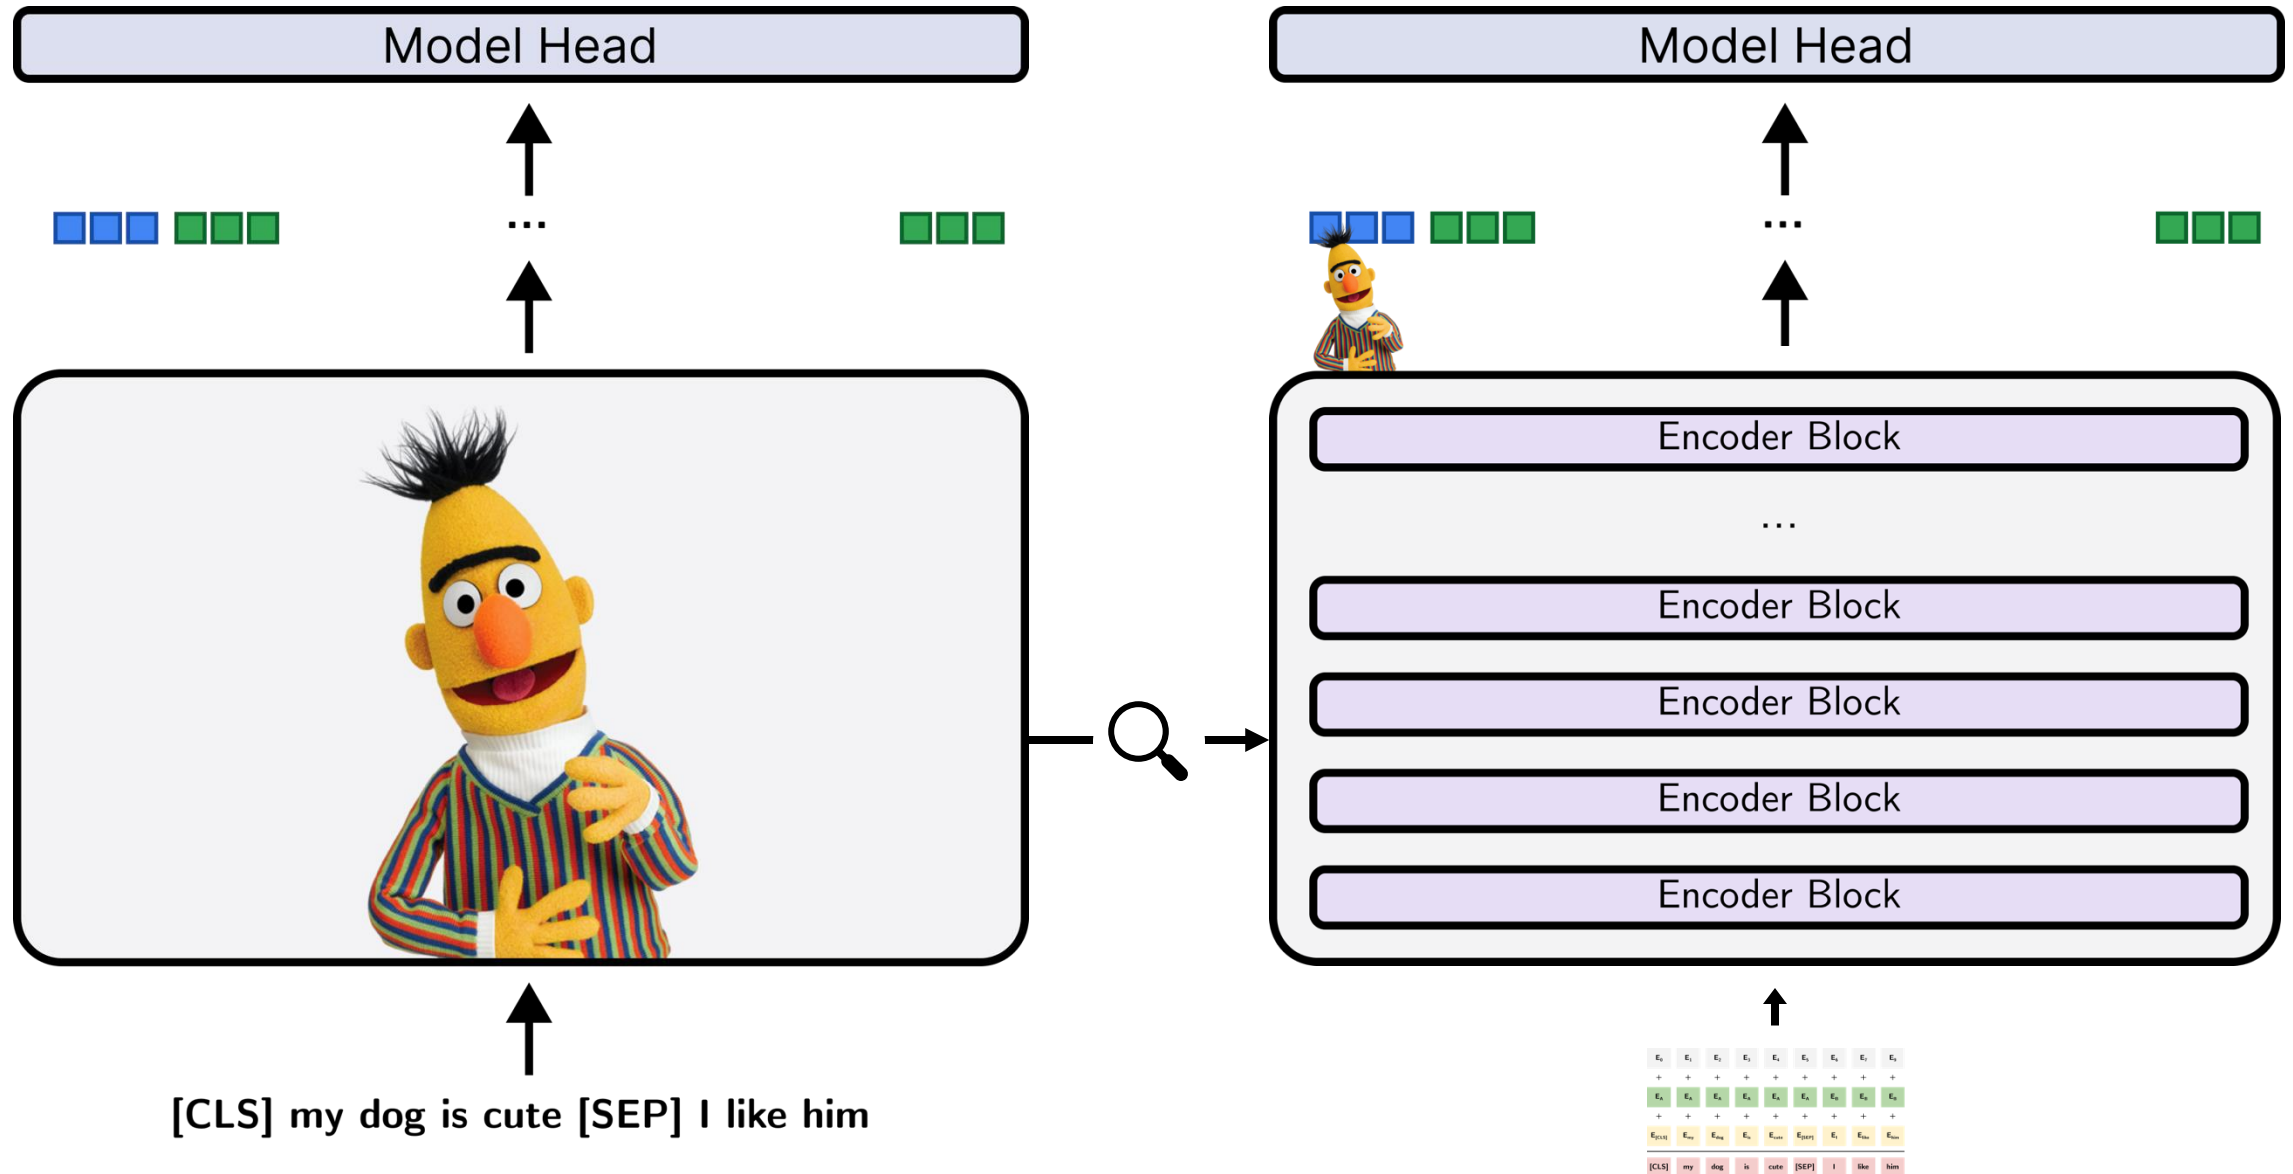
A bird's eye view of the core architecture

The input sequence is projected input an embedding space before adding token-, segment- and positional embeddings.



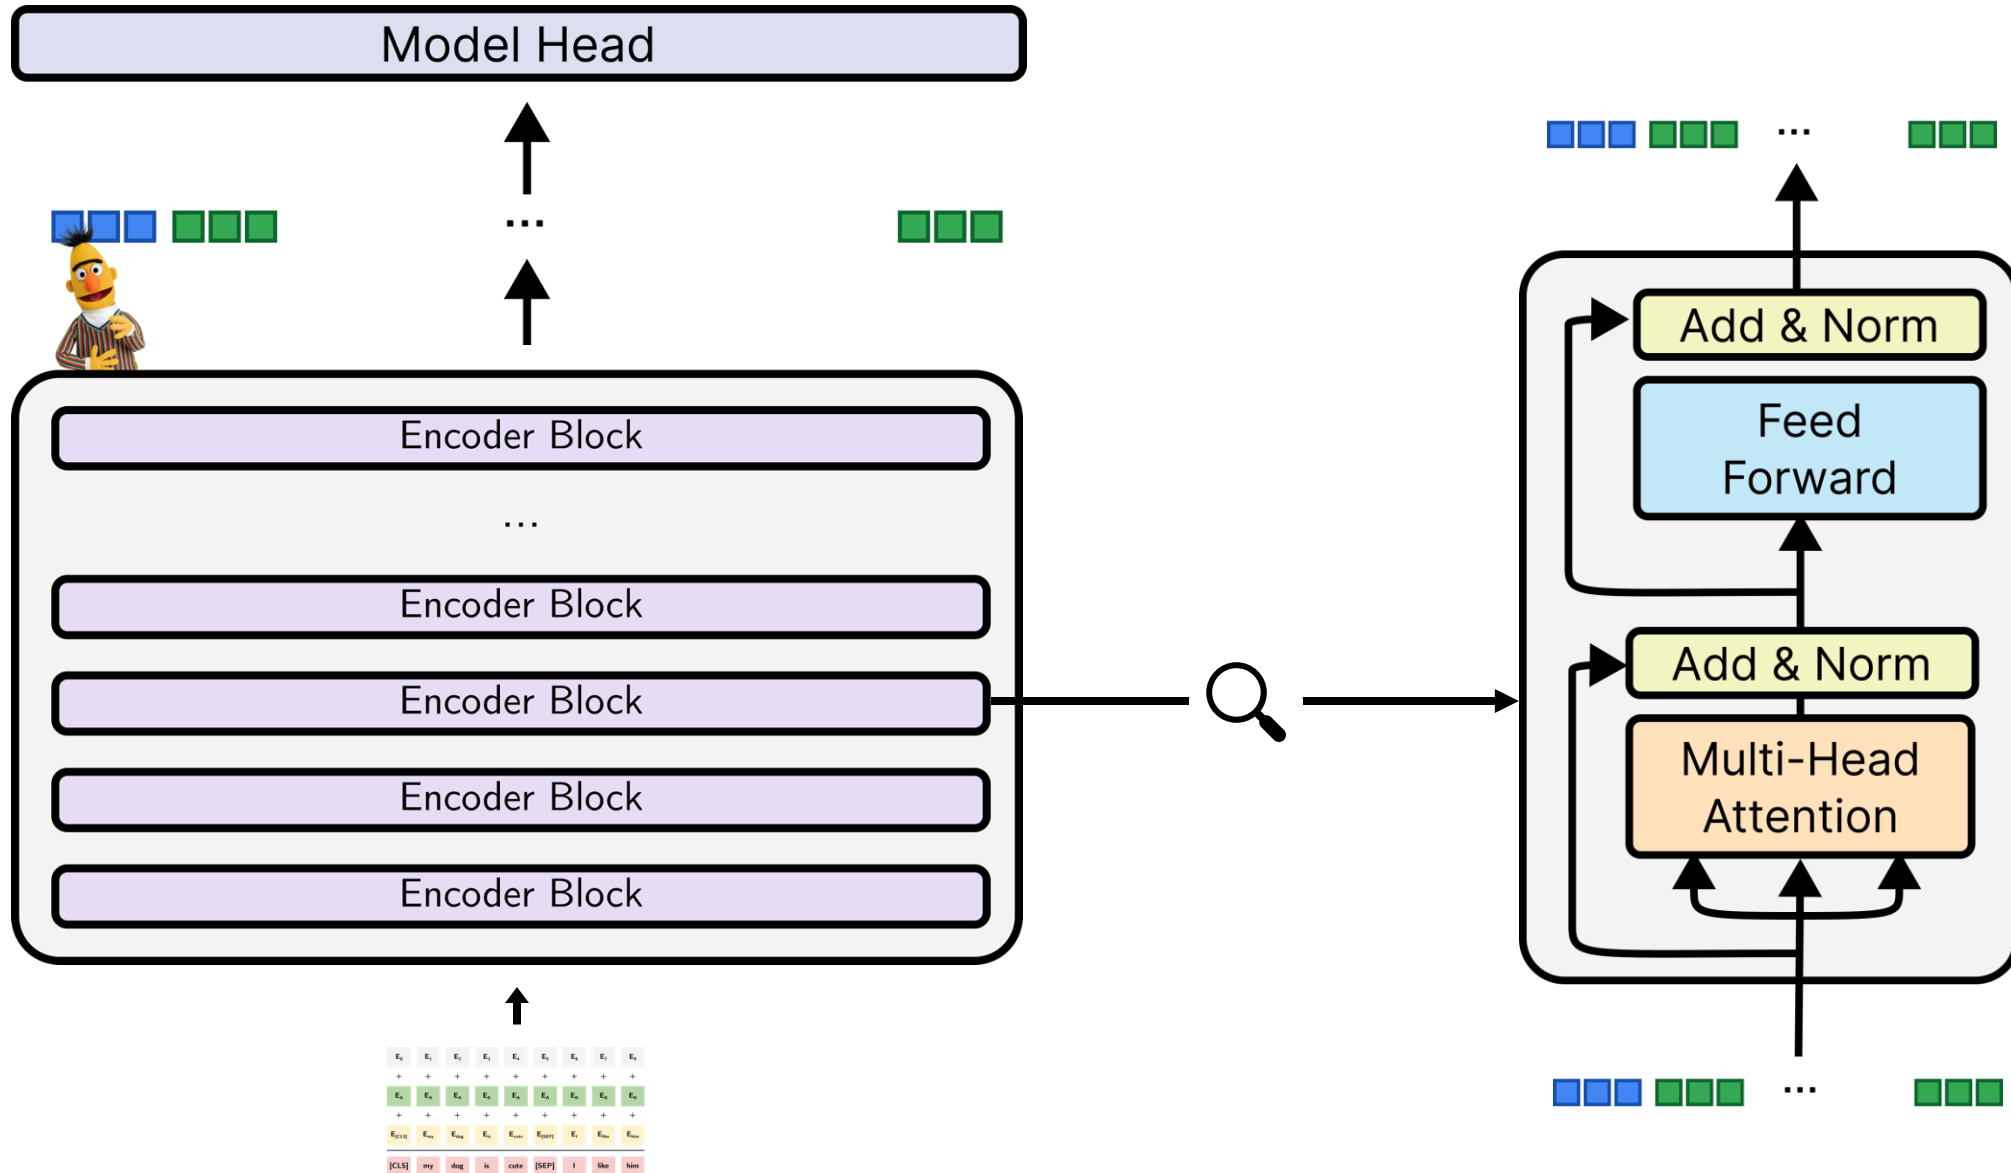
A bird's eye view of the core architecture

BERT's model body is a repetition of multiple sequentially concatenated equivalent encoder blocks.



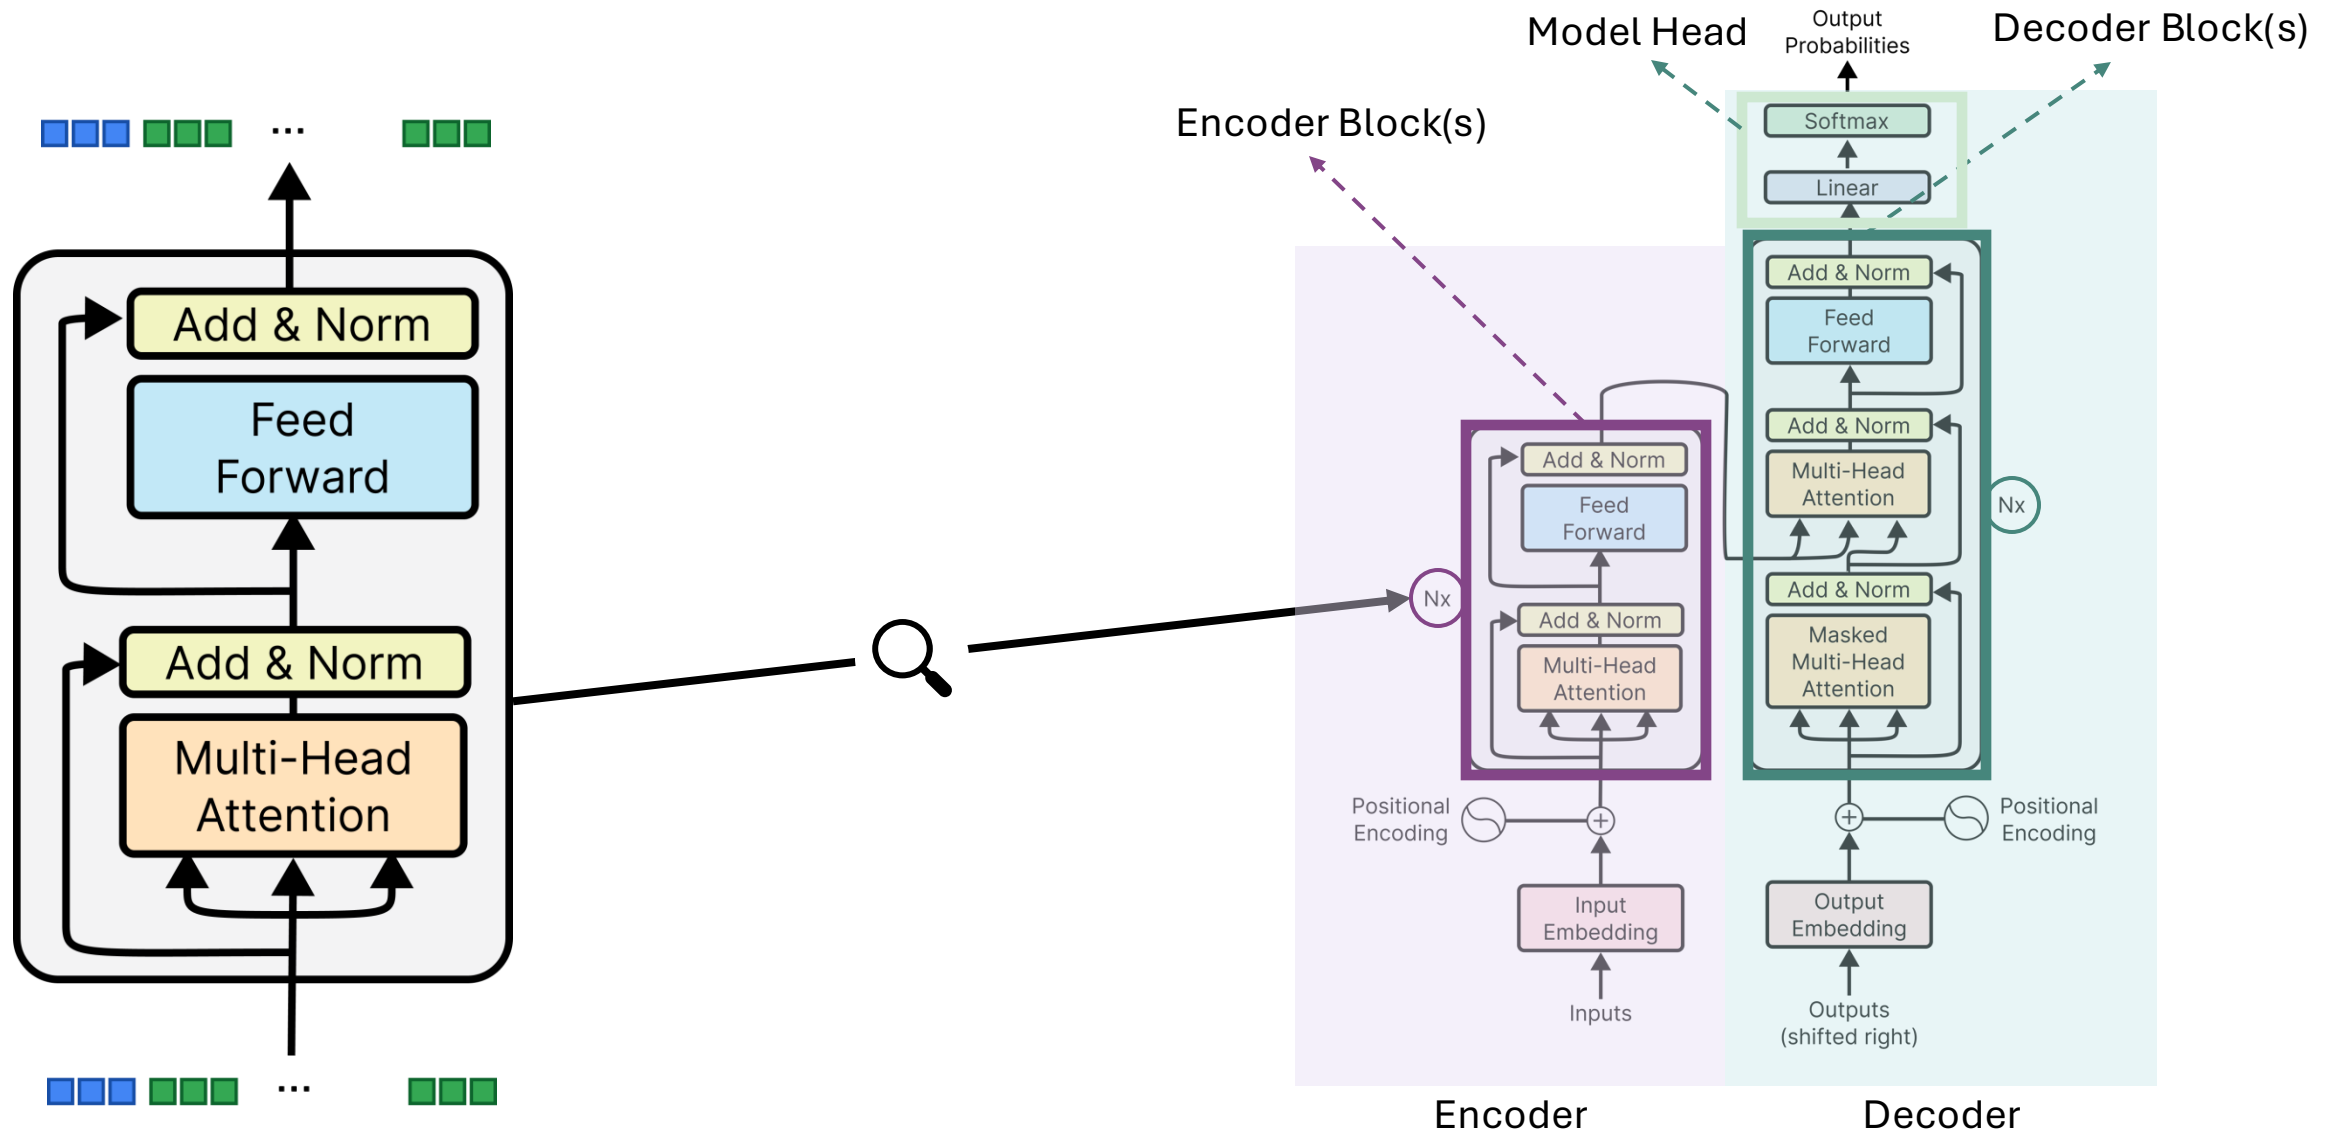
A bird's eye view of the core architecture

The encoder block consists of multi-head attention and a feed forward network with layer normalization and residual connectoins.



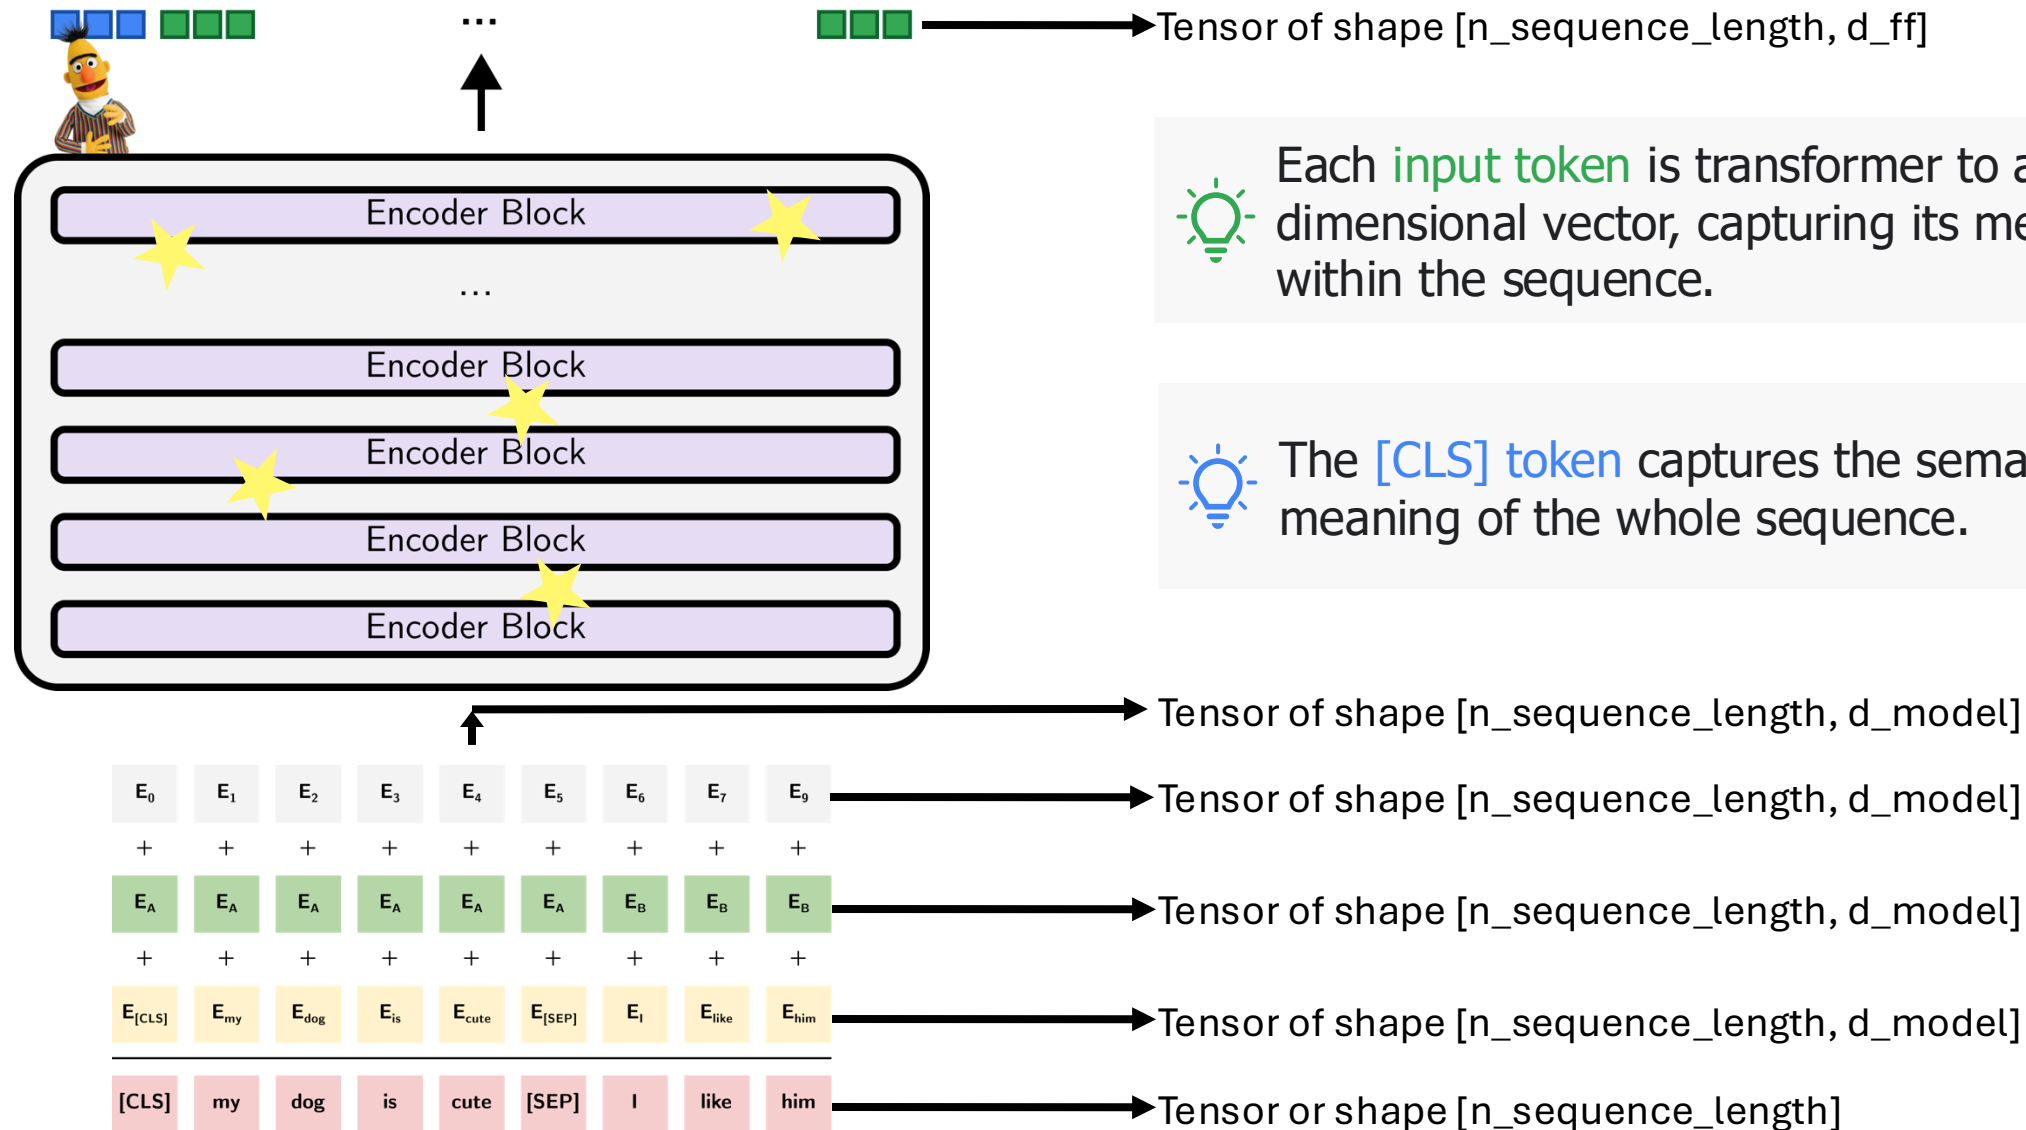
A bird's eye view of the core architecture

BERT “steals” its encoder block from the transformer architecture.



A bird's eye view of the core architecture

BERT transforms a token-sequence into a sequence of dense vectors capturing their contextual meanings.



d_ff stands for the output dimension of the feed forward network of the encoder block.

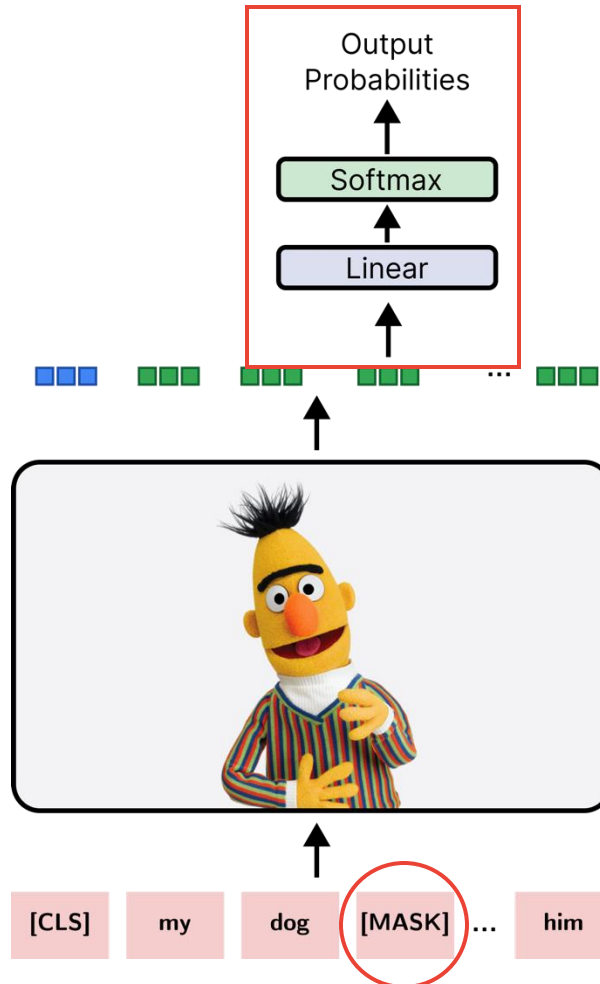
BERT's versatility

The pre-training / fine-tuning process

BERT's versatility - The pre-training / fine-tuning process

BERT learns contextual language modelling by partially masking out input tokens.

🎯 Mask a 15% of the input at random and predict those masked tokens.



› 80% of the time: Replace the word with the *[MASK]* token
To learn how words relate to their surrounding context to make accurate predictions.



› 10% of the time: Replace the token with a random token
To prevent the model from simply learning that it should pay special attention to *[MASK]* tokens.

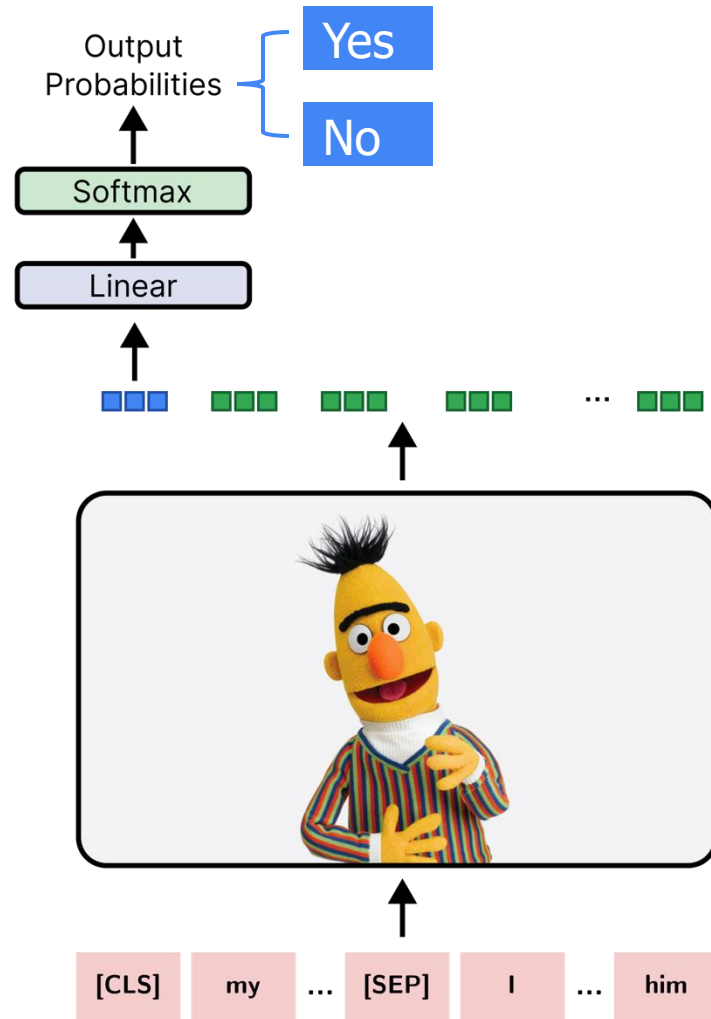


› 10% of the time: Keep the word unchanged
To bias the representation towards the actual observed word.

BERT's versatility - The pre-training / fine-tuning process

BERT learns sentence level contextual language modeling through next sentence predictions.

🎯 To learn in order sentence relationships, we pre-train for a binarized next sentence prediction.

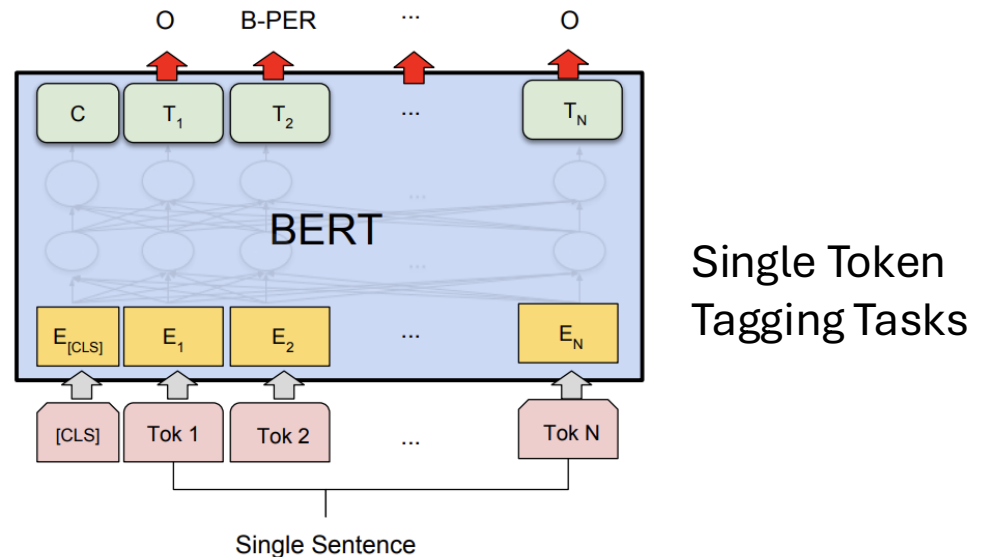
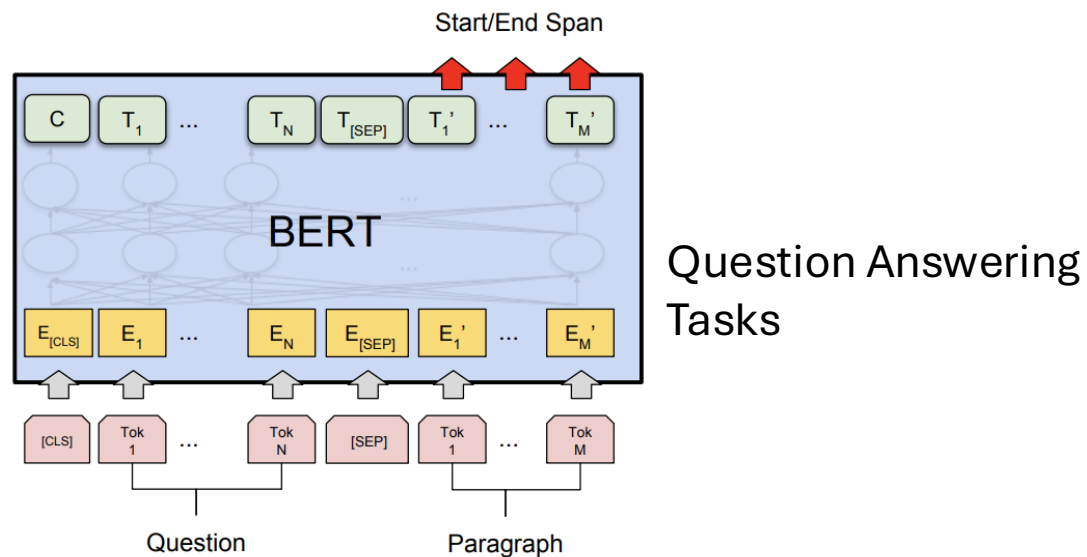
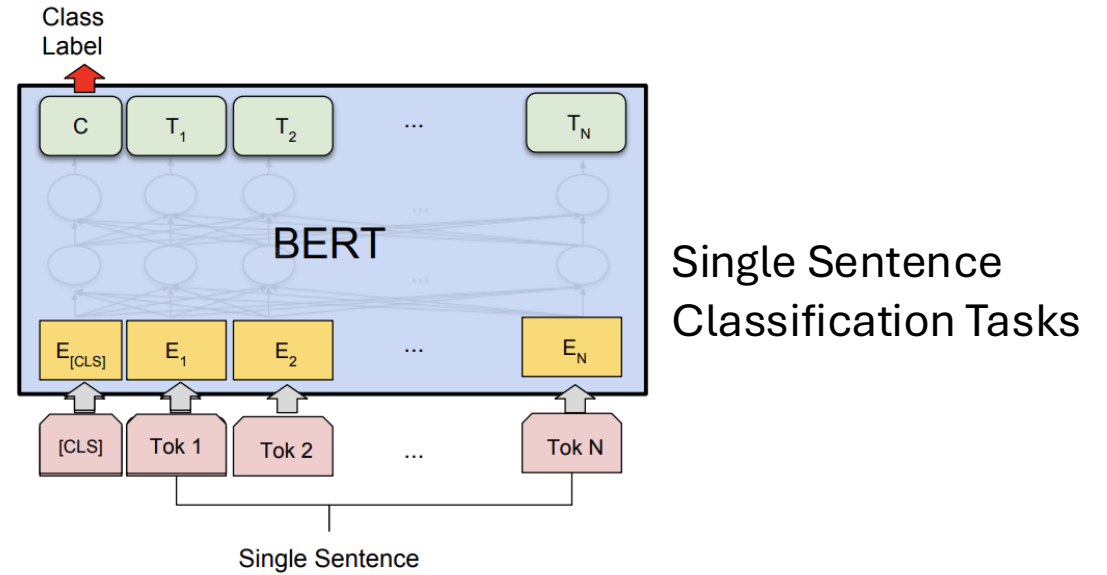
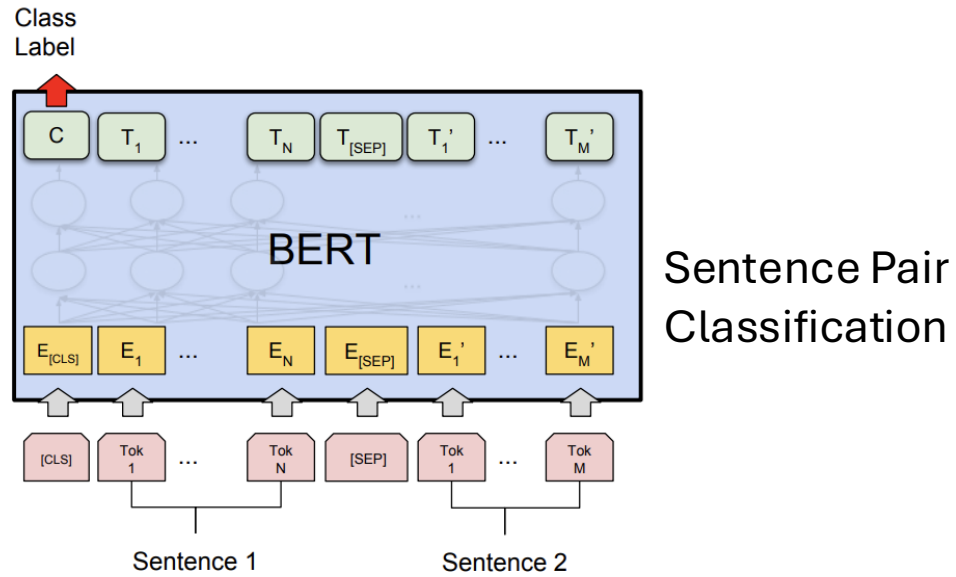


› 50% of the time actual next sentence is used 50% of the time it is a random sentence from the corpus

› [CLS] related output is used for binary classification

BERT's versatility - The pre-training / fine-tuning process

We can adapt BERT to different tasks by using different model heads and fine-tuning all the parameters end-to-end.



Conclusion & Discussion

And a few numbers



Conclusion & Discussion

Some numbers

	Encoder Blocks	$d_{\text{Feed Forward \& Embeddings}}$	Attention Heads	Total Params	Pretraining	TPUs
BERT _{BASE}	12	768	12	110 millions	4 days	4
BERT _{LARGE}	24	1024	16	340 millions	4 days	16*

* At the current on-demand prices of v3 TPUs - which were most likely used at that time – of 2.2\$ the large model would cost around 10k of pretraining. Keep in mind that Google trained it on their own infrastructure, which will have cost them much less.

Dataset

- BooksCorpus (800 million words) & text passages of english wikipedia (2,5 billion words) with WordPiece tokenization with 30,000 token vocabulary

Batching/Training Config

- Total sequence length ≤ 512 tokens¹
- Batch size of 256 sequences (256 sequences x 512 tokens = 128,000 tokens/batch)
- 1,000,000 total steps, which is approximately 40 epochs over the 3.3 billion word corpus.

For the geeks

- Adam optimizer with lr of $1e-4$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, L2 weight decay of 0.01, warmup over the first 10,000 steps, and linear lr decay
- Dropout of 0.1

1. To speed up pretraing a sequence length of 128 for 90% of the steps was used and only the final 10% of the steps of sequence of 512 to learn the positional embeddings.

Conclusion & Discussion

GLUE was used as the primary benchmark for BERT which it dominated in 2019.

Back in 2019

System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4
OpenAI GPT	82.1/81.4	70.3	87.4	91.3
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9

System	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average -
Pre-OpenAI SOTA	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	60.5	86.5	89.3	70.1	82.1

And Now?

Rank	Name	Model	Score	CoLA	SST-2	MRPC	STS-B	QQP	MNLI-m	MNLI-mm	QNLI	RTE	WNLI	AX
1	Microsoft Alexander v-team	Turing ULR v6	91.3	73.3	97.5	94.2/92.3	93.5/93.1	76.4/90.9	92.5	92.1	96.7	93.6	97.9	55.4
2	JDExplore d-team	Vega v1	91.3	73.8	97.9	94.5/92.6	93.5/93.1	76.7/91.1	92.1	91.9	96.7	92.4	97.9	51.4
3	Microsoft Alexander v-team	Turing NLR v5	91.2	72.6	97.6	93.8/91.7	93.7/93.3	76.4/91.1	92.6	92.4	97.9	94.1	95.9	57.0
4	DIRL Team	DeBERTa + CLEVER	91.1	74.7	97.6	93.3/91.1	93.4/93.1	76.5/91.0	92.1	91.8	96.7	93.2	96.6	53.3
...														
47	Mikita Sazanovich	Routed BERTs	80.7	56.1	93.6	88.6/84.7	88.0/87.6	71.0/88.8	85.2	84.5	92.6	80.0	65.1	9.2
48	USCD-AI4Health Team	CERT	80.7	58.9	94.6	89.8/85.9	87.9/86.8	72.5/90.3	87.2	86.4	93.0	71.2	65.1	39.6
49	Jacob Devlin	BERT: 24-layers, 16-heads, 1024-hidden	80.5	60.5	94.9	89.3/85.4	87.6/86.5	72.1/89.3	86.7	85.9	92.7	70.1	65.1	39.6
50	Chen Qian	KerasNLP XLM-R	80.4	56.3	96.1	89.8/86.3	88.4/87.7	72.3/89.0	87.7	87.1	92.8	69.2	65.1	40.6

General **L**anguage **U**nderstanding **E**valuation (GLUE) is a collection of diverse natural language understanding tasks.

Conclusion & Discussion

An overview of the datasets used to evaluate BERT.

Acronym	Full Name	Description
MNLI	Multi-Genre Natural Language Inference	Predicting whether sentence pairs are <i>entailment</i> , <i>contradiction</i> , or <i>neutral</i>
QQP	Quora Question Pairs	Determining if two questions are semantically equivalent
QNLI	Question Natural Language Inference	A version of Stanford Question Answering Dataset converted to binary classification
SST-2	Stanford Sentiment Treebank	Binary sentiment classification of movie reviews
CoLA	Corpus of Linguistic Acceptability	Predicting whether English sentences are linguistically acceptable
STS-B	Semantic Textual Similarity Benchmark	Rating similarity of sentence pairs on a scale of 1-5
MRPC	Microsoft Research Paraphrase Corpus	Identifying whether sentence pairs are semantically equivalent
RTE	Recognizing Textual Entailment	Similar to MNLI but with much less training data
WNLI	Winograd NL	A small natural language inference dataset ¹

1. Authors excluded this due to issues with the dataset)

Conclusion & Discussion

One can extract fixed features from the pretrained model.

System	Dev F1	Test F1
ELMo (Peters et al., 2018a)	95.7	92.2
CVT (Clark et al., 2018)	-	92.6
CSE (Akbik et al., 2018)	-	93.1
Fine-tuning approach		
BERT _{LARGE}	96.6	92.8
BERT _{BASE}	96.4	92.4
Feature-based approach (BERT _{BASE})		
Embeddings	91.0	-
Second-to-Last Hidden	95.6	-
Last Hidden	94.9	-
Weighted Sum Last Four Hidden	95.9	-
Concat Last Four Hidden	96.1	-
Weighted Sum All 12 Layers	95.5	-

CoNLL-2003 Named Entity Recognition results. Hyperparameters were selected using the Dev set. The reported Dev and Test scores are averaged over 5 random restarts using those hyperparameters.

1. P ... Precision, R ... Recall, TP ... True Positives, FP ... False Positives, FN ... False Negatives

2. This represents the personal opinion of the presenter. He does not have an affiliation with BERT or Google to praise or advise against any models/products.

Reminder¹

- $F1 = 2 * (P * R) / (P + R)$
- $P = TP / (TP + FP)$
"Of all entities we predicted, what fraction did we get right?"
- $R = TP / (TP + FN)$
"Of all actual entities, what fraction did we find?"

 Word2Vec²

- › Why would we like to do that?

Authors: Computational Efficiency & adaptability to task specific model architectures.

- › Why might it make sense to use different layers instead of just the final layer?


A few guesses from the presenter:

- *Different layers capture different types of information*
- *Combining layers might provide complementary information*
- *Final layer might be specialized on pretraining task*




Conclusion & Discussion


Can you name other models inspired by the Sesame Street?




ELMO (1980)



GROVER (1970)




ELMo (2018)




GROVER (2019)

Embeddings from Language Models

Generates realistic and controlled fake news



ERNIE (1969)



ERNIE (2019)

BERT with entity-level and phrase-level masking

Conclusion & Discussion

The BERT architecture has a huge ecosystem¹ of different specialised BERT-like models

- ALBERT
- RoBERTa
- HerBERT
- [...]
- I-BERT

❓ You-BERT?

Have you ever used a BERT-like model? What for / What was your experience?

❓ Isn't a decoder all you need?

Will BERT like models be swallowed by decoder only models (e.g. GPT)



That's it!

1. https://huggingface.co/docs/transformers/model_doc/bert

APENDIX

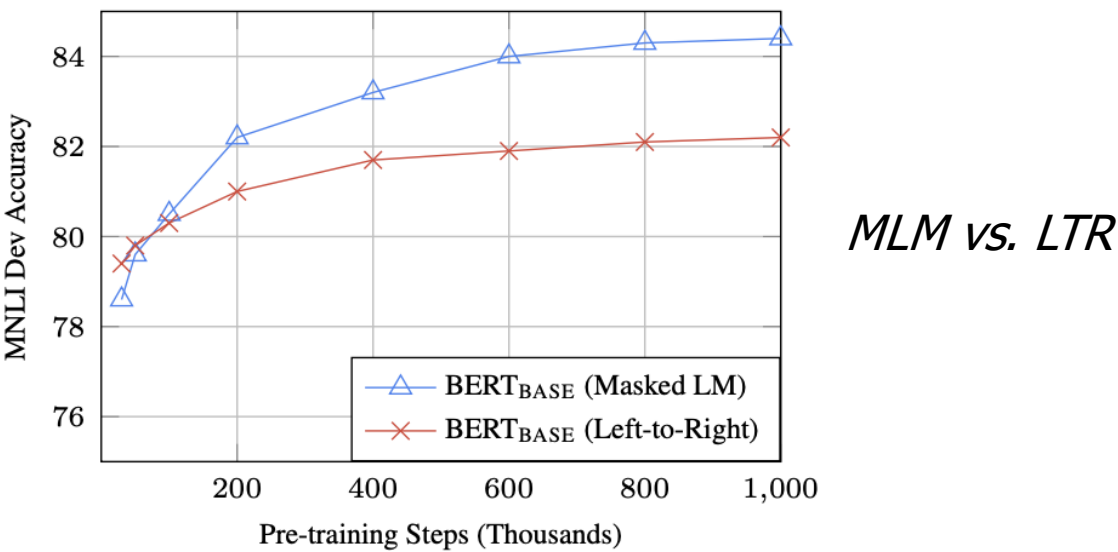
Ablation Studies

The effect of the different train strategies on the overall performance of BERT.

Tasks	Dev Set				
	MNLI-m (Acc)	QNLI (Acc)	MRPC (Acc)	SST-2 (Acc)	SQuAD (F1)
BERT _{BASE}	84.4	88.4	86.7	92.7	88.5
No NSP	83.9	84.9	86.5	92.6	87.9
LTR & No NSP	82.1	84.3	77.5	92.1	77.8
+ BiLSTM	82.1	84.1	75.7	91.6	84.9

Table 5: Ablation over the pre-training tasks using the BERT_{BASE} architecture. “No NSP” is trained without the next sentence prediction task. “LTR & No NSP” is trained as a left-to-right LM without the next sentence prediction, like OpenAI GPT. “+ BiLSTM” adds a randomly initialized BiLSTM on top of the “LTR + No NSP” model during fine-tuning.

Note: I have not found any indication on how the model is modified for LTR (i.e. whether it’s only a change in the loss function or whether they’ve added future token masking like in decoder only models)

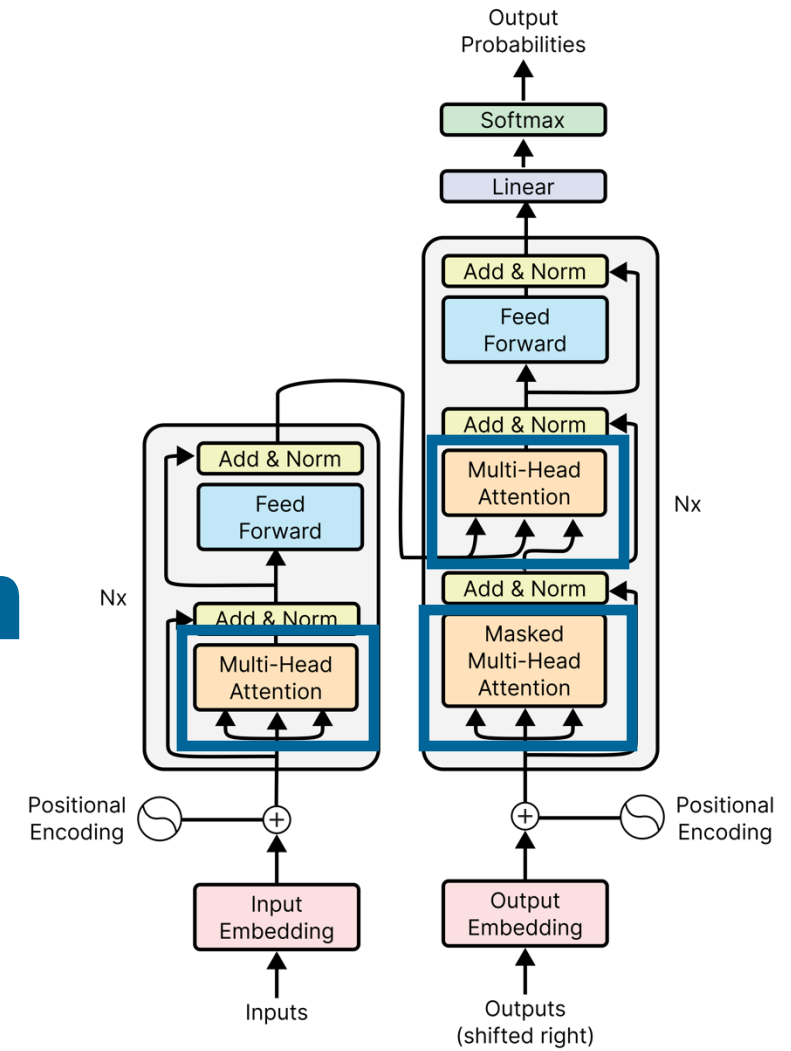


Masking Rates			Dev Set Results		
MASK	SAME	RND	MNLI	NER	
			Fine-tune	Fine-tune	Feature-based
80%	10%	10%	84.2	95.4	94.9
100%	0%	0%	84.3	94.9	94.0
80%	0%	20%	84.1	95.2	94.6
80%	20%	0%	84.4	95.2	94.7
0%	20%	80%	83.7	94.8	94.6
0%	0%	100%	83.6	94.9	94.6

Performance of different masking strategies

The attention mechanism

The cornerstone of the transformer's ability to capture context.



The attention mechanism

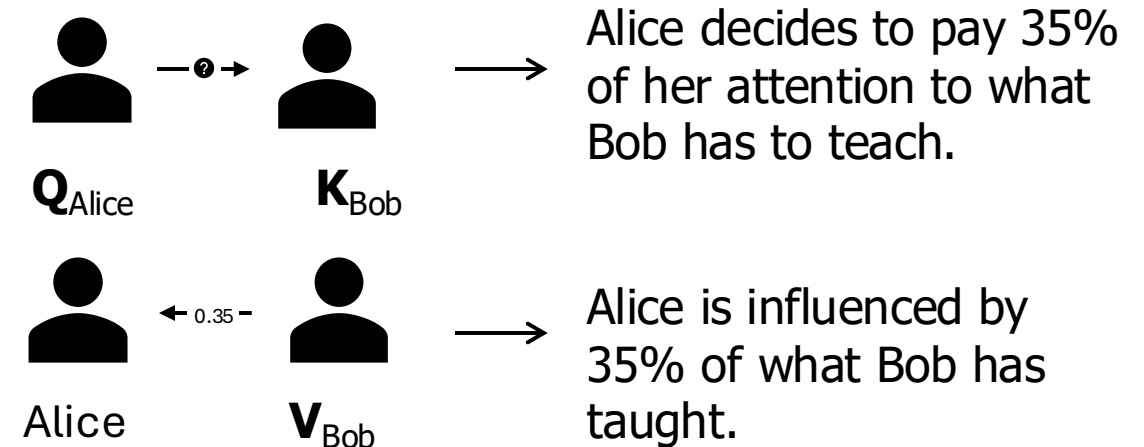
An example of the exchange of attention in a group of people eager to share and gather knowledge.

- Alice, Bob, Charlie, David, and Eve want to advance their knowledge in certain areas by spending time and attention.
- Each of them has some knowledge they can teach.
- Each of them has a **limited capacity** to **pay attention** (one can only learn so much in a week).
- Each of them can **receive attention indefinitely** (you can be listened to by everyone).


Query ... a description of the knowledge they want to gather

Key ... a description of what they can teach

Value ... the actual knowledge



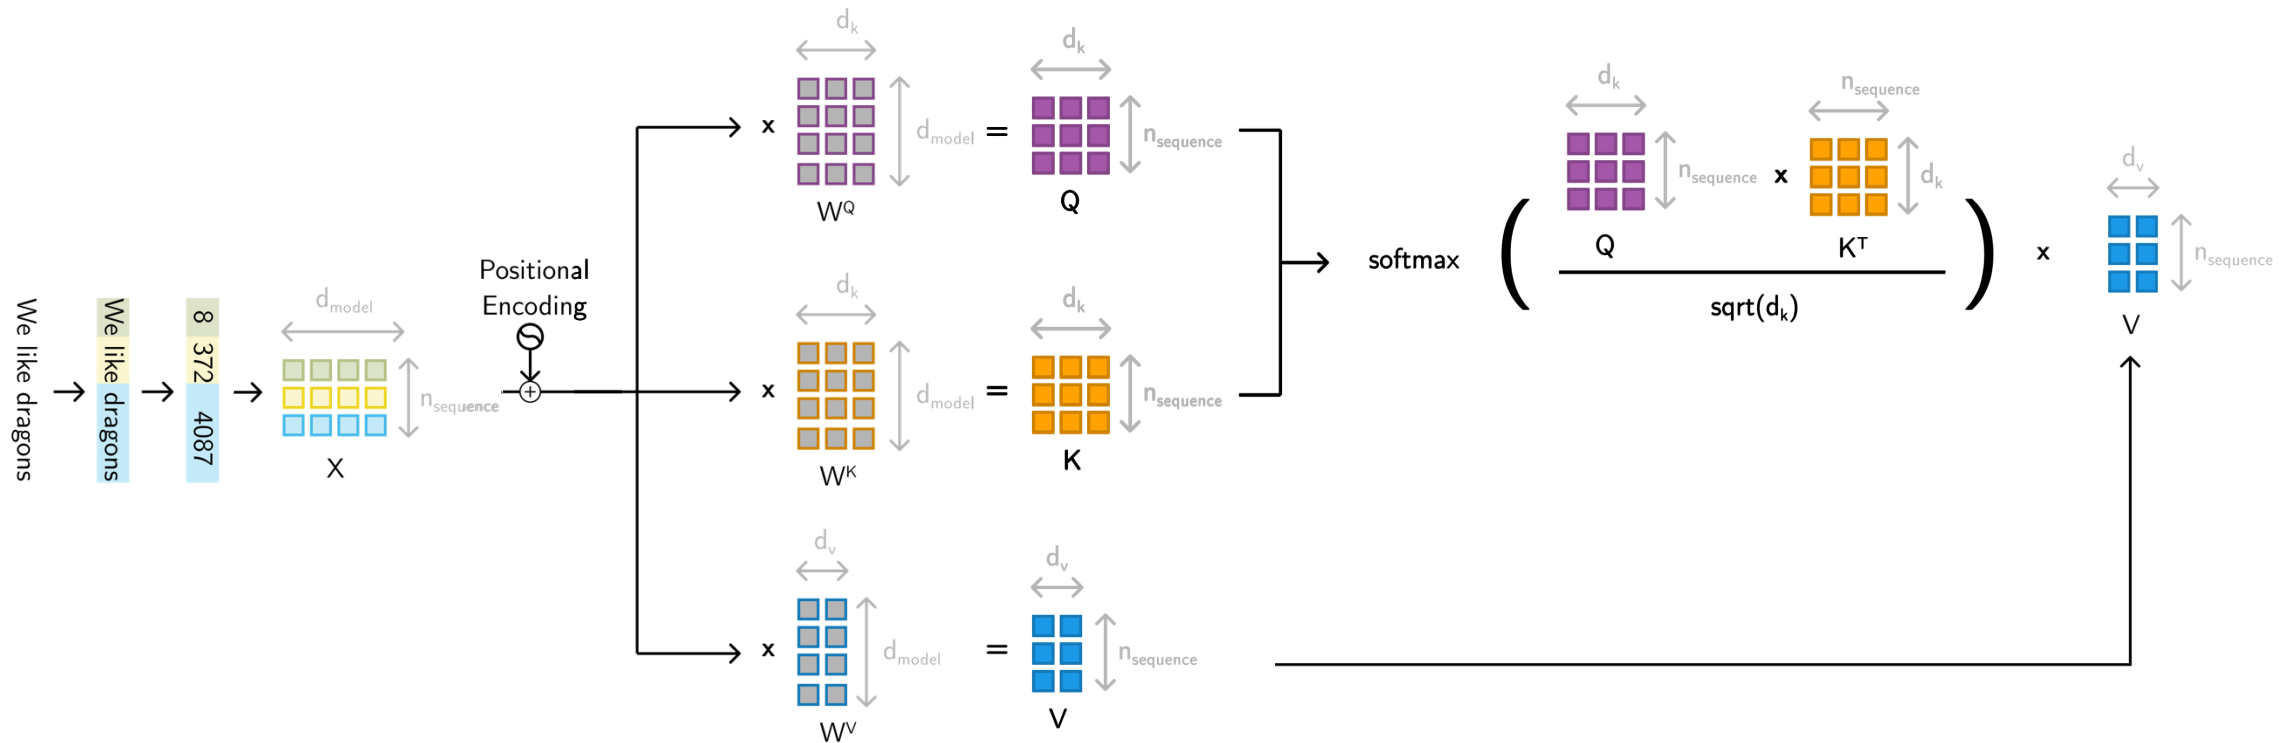
Visualisation of an interaction between two people¹

 The goal of the attention mechanism is to compute and ingest *how much* and *how* each token should influence the representation ("meaning") of every other token within the input sequence.

1. Corresponds to two tokens

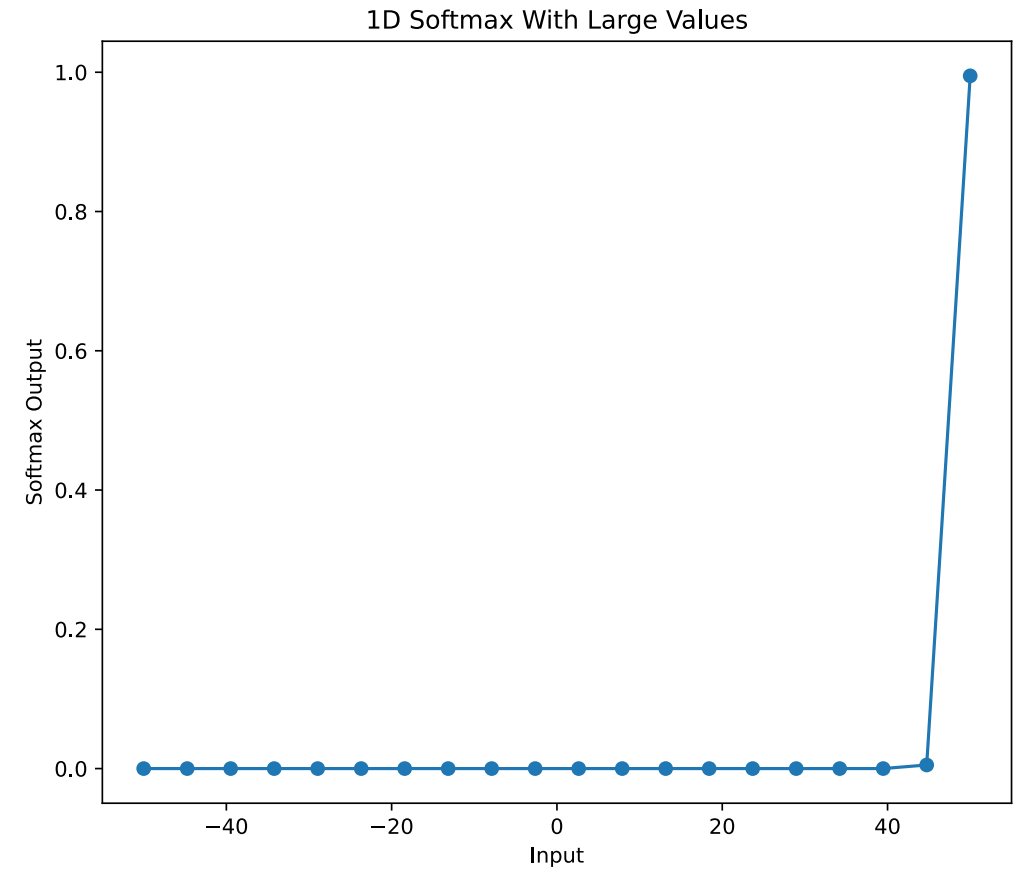
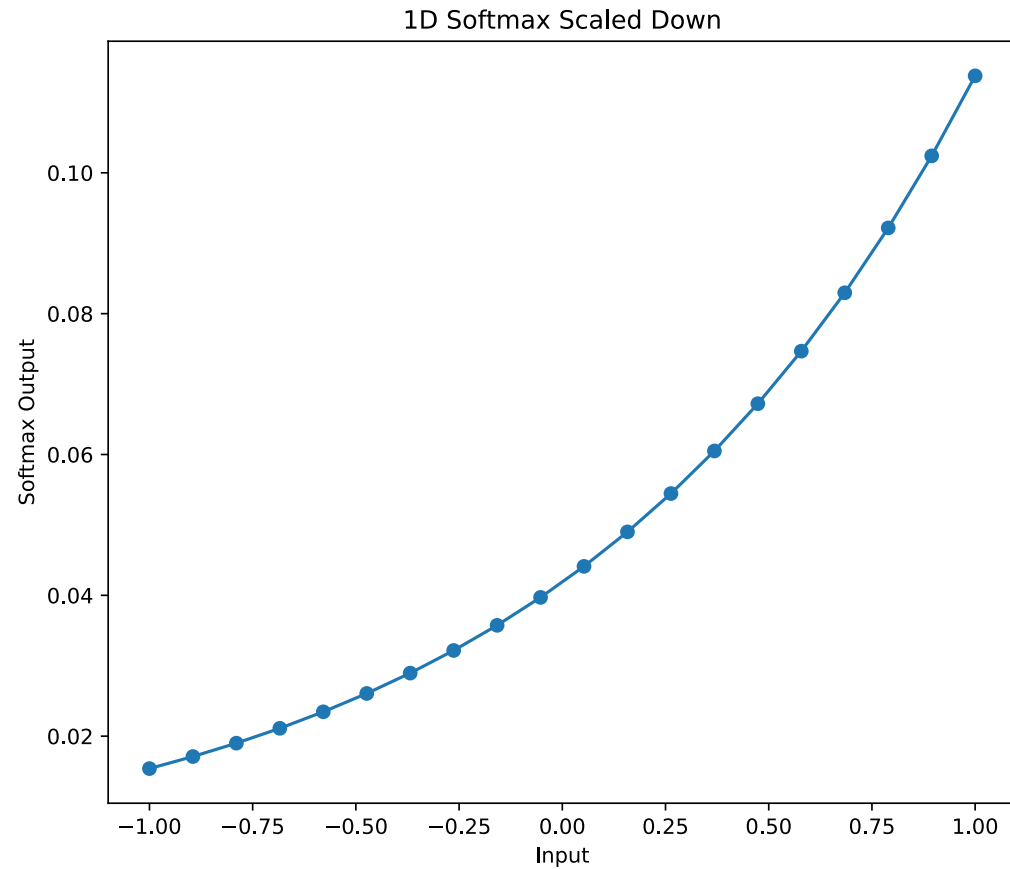
The attention mechanism

Compute and ingest *how much* and *how* each token should influence the representation of every other token.




The attention mechanism

Scaling the softmax input smooths the curve and ensures gradient stability.



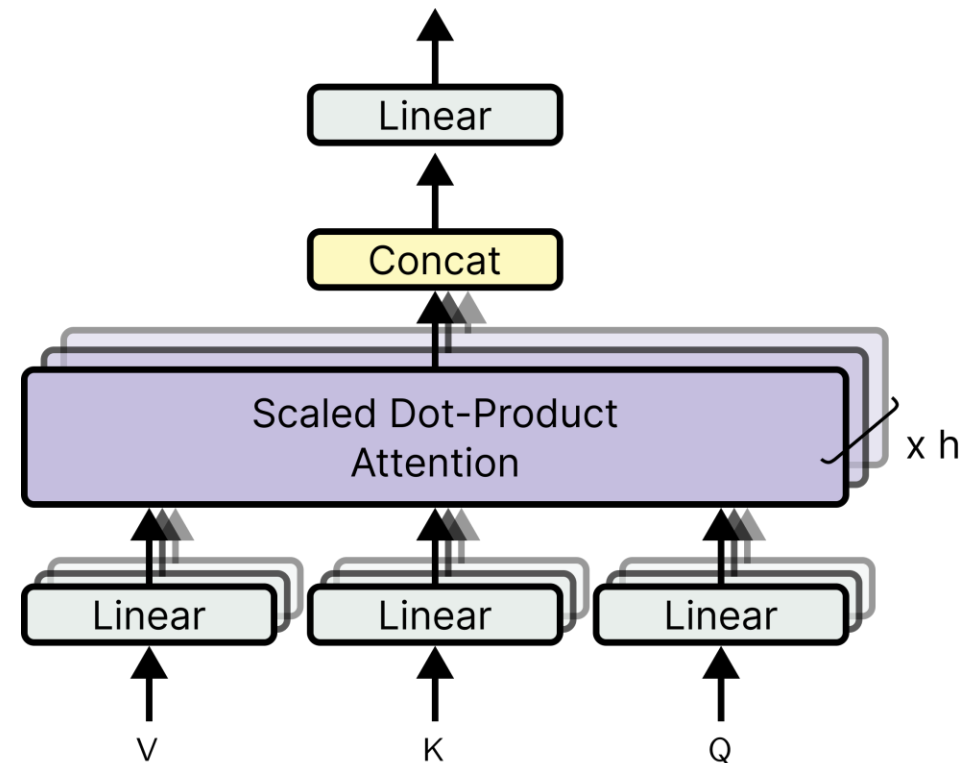
The attention mechanism

Example of randomly assigned attention scores to demonstrate the ability of the softmax function to normalise its input.




Multi-Head Attention


How transformers pay attention to different "semantic properties" of tokens



The attention mechanism

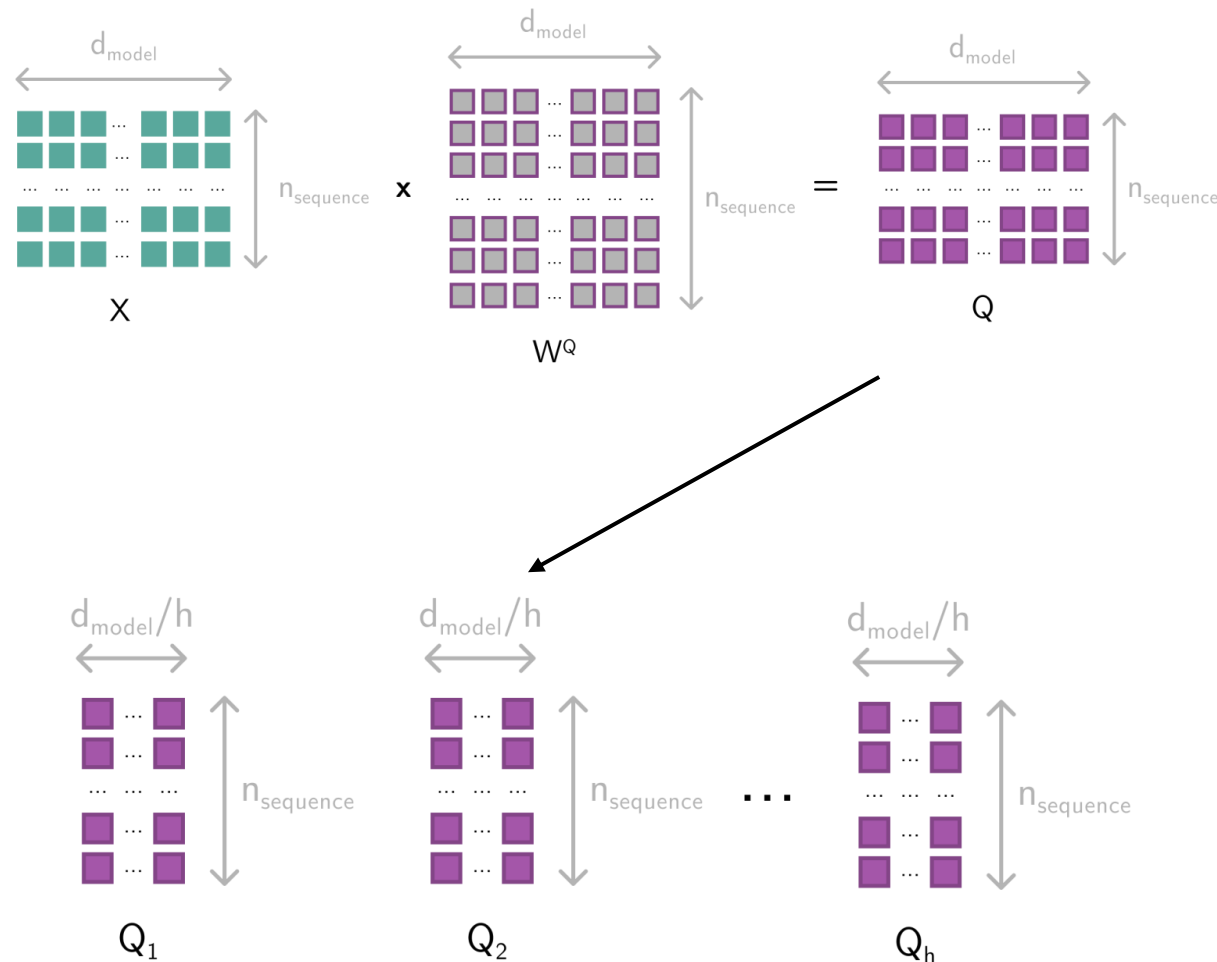
Mathematical intuition behind multi-head attention.

 The goal of the multi-head attention is to attend to tokens from different perspectives in parallel.



The attention mechanism

Mathematical intuition behind multi-head attention.



```
import torch
from torch import nn
```

```
d_model = 256
n_context = 1048
h = 8
d_k = d_model // h
```

```
x = torch.randn(n_seq, d_model)
```

```
W_Q = nn.Linear(d_model, d_model, bias=False)
W_K = nn.Linear(d_model, d_model, bias=False)
W_V = nn.Linear(d_model, d_model, bias=False)
W_O = nn.Linear(d_model, d_model, bias=False)
```

```
Q = W_Q(x)
K = W_K(x)
V = W_V(x)
```

```
Q = Q.view(n_seq, h, d_k).transpose(0, 1)
K = K.view(n_seq, h, d_k).transpose(0, 1)
V = V.view(n_seq, h, d_k).transpose(0, 1)
```

💡 In multi-head attention, we produce h slices each of shape n_{seq}, d_k and stack these matrices along the first dimension. Q, K and V have the shape h, n_{seq}, d_k

The attention mechanism

Mathematical intuition behind multi-head attention.

```
import torch
from torch import nn

d_model = 256
n_seq = 1048
h = 8
d_k = d_model // h

x = torch.randn(n_seq, d_model)

W_Q = nn.Linear(d_model, d_model, bias=False)
W_K = nn.Linear(d_model, d_model, bias=False)
W_V = nn.Linear(d_model, d_model, bias=False)
W_O = nn.Linear(d_model, d_model, bias=False)

Q = W_Q(x) # [n_seq, d_model]
K = W_K(x) # [n_seq, d_model]
V = W_V(x) # [n_seq, d_model]

# [h, n_seq, d_k]
Q = Q.view(n_seq, h, d_k).transpose(0, 1)
K = K.view(n_seq, h, d_k).transpose(0, 1)
V = V.view(n_seq, h, d_k).transpose(0, 1)

K_T = K.transpose(-2, -1)
```

```
attn_scores = torch.matmul(Q, K_T) / torch.sqrt(torch.tensor(d_k))
attn_weights = nn.functional.softmax(attn_scores, dim=-1)
attn = torch.matmul(attn_weights, V) # [h, n_seq, d_k]

attn = attn.transpose(0, 1).contiguous().view(n_seq, d_model)

output = W_O(self_attention)
```

$$MultiHead(Q, K, V) = Concat(head_1, head_2, \dots, head_h)W^O$$

$$head_i = Attention(Q_i, K_i, V_i)$$




Multi-head attention learns contextual dependencies from different perspectives in parallel by splitting the **Q**uery, **K**ey and Value matrices into multiple heads.

contiguous(): Ensures the tensor is stored in contiguous memory; doesn't change the shape, but prepares for the view operation


Batching

Simultaneously process multiple training samples



Batching

Batching allows us to train on multiple samples in parallel. All sequences need to be padded to the same length.



Note: d_{model} , n_{sequence} and n_{batch} are usually significantly larger; however it generally holds $n_{\text{batch}} < d_{\text{model}} < n_{\text{sequence}}$